

MaceStyle: Bringing It In-House

Running our Style Guide validator on Mace's own billing and governance

Prepared for **Tobias** by Stephen Cummins · AI Service Lead, Mace Digital

2 July 2026 · updated 5 July 2026

MaceStyle — the automated Style Guide validator already piloting on the Mace Way Control Centre — now runs its AI on **GPT via Azure OpenAI: Microsoft-native**, on Azure billing and governance, at **pennies per document**. I've already made the switch and proven it end-to-end on the live pilot. To make it a Mace service it needs a Mace Azure subscription to own it and a data-governance nod. Then it's roughly half a day to cut over.

What we have today

MaceStyle reads documents in the Mace Way Control Centre and checks them against the Writing Style Guide — British spelling, contractions, symbols, number and font standards — flagging issues and proposing corrections. It is **live in pilot**, validated by testers (Natasha, Jade) in June. The intelligent rules — the judgement calls a simple find-and-replace can't make — are powered by AI.

What changed this month: I re-architected MaceStyle so the AI backend is *swappable* via a single setting, and switched the live pilot onto **GPT-5 (Azure OpenAI)** — Microsoft's own AI service, running inside Azure. It's proven working end-to-end. The same code can run Claude instead with no rewrite, so we are not locked to one model.

The catch: it currently runs on *my personal* Azure subscription. That proves the concept — it's not how Mace should run a service it depends on. To make this a Mace service, the billing and governance need to sit inside Mace.


The proposal: run it on Mace's Azure — Microsoft-native

Mace's instinct is to stay on Microsoft wherever possible. This fits that exactly. The AI now runs on **Azure OpenAI** — the GPT models (the ChatGPT family) delivered as a first-party Microsoft Azure service:

	TODAY (MY AZURE)	PROPOSED (MACE'S AZURE)
Who pays	Me, personally	Mace's existing Azure invoice
What it is	GPT via Azure OpenAI	Same — a native Azure service Mace already has
Governance	Mine	Entra ID, RBAC, data zone, spend caps
Procurement	—	None. No new vendor, no third-party sign-up

The key point: **there is no new supplier and no AI contract to negotiate**. Azure OpenAI is Microsoft, on the Azure agreement Mace already runs. Standing it up in a Mace subscription is deploying a model and pointing the app at it — no marketplace step, no external account.

Model choice stays open. Because the backend is swappable, Mace can run **Claude** instead (via Microsoft Foundry, also inside Azure) for content that benefits from its slightly more careful editing — a one-setting change, not a rebuild. My recommendation: **default to GPT / Azure OpenAI** (easiest to provision and the cheaper of the two) and keep Claude available as an option.



How MaceStyle runs under Mace: the Function calls GPT via Azure OpenAI, all inside Mace's Azure. Claude (via Foundry) is a one-setting alternative.

What it costs

Consumption-based and genuinely cheap. I measured both models on real documents through the live pipeline:

DOCUMENT SIZE	GPT-5-MINI (RECOMMENDED)	CLAUDE (ALTERNATIVE)
~1 page	£0.003	£0.004
~3 pages	£0.004	£0.008
~8 pages	£0.008	£0.015

US-dollar list prices, converted approximately; Mace's enterprise Azure agreement would likely discount further.

GPT is **~35–50% cheaper per document** than Claude, and the gap widens with length — but in absolute terms both are a **rounding error**: realistic Control Centre volume is a few pounds a month. No upfront commitment; Mace pays only for what it uses. Cost isn't the deciding factor — ease of ownership is, and Azure OpenAI wins on both.

What I need from you

1. **A home for it** — nominate a Mace Azure subscription and cost centre to own the service. Mace already runs on Azure (the *mace365* tenant, Power Platform), so nothing new is stood up.
2. **A governance nod** — document text is processed by a US-hosted model. Azure OpenAI's data zone and Azure-native governance are built for exactly this; it needs Sapna's sign-off, which I'll tee up.
3. **Half a day of my time** to cut over — I've already built and proven the whole path on the live pilot, so this is execution, not discovery.

■ Why this matters beyond MaceStyle

MaceStyle is the **proof, not the point**. It shows a repeatable pattern: identify a real Mace problem, stand up an AI service to solve it, and hand it over running on Mace's own Microsoft/Azure billing and governance — safely, cheaply, and without a procurement cycle.

That pattern is the AI Service Lead role in practice. Give me the mandate and MaceStyle becomes the **first of several** — each one an internal AI service that pays for itself in hours saved and lands inside Mace's existing Azure controls from day one.

Technical detail for whoever provisions it: see the companion *Handover Runbook* (covers both the Azure OpenAI / GPT and Foundry / Claude setups). · Stephen Cummins · AI Service Lead, Mace Digital · stephen.cummins@macegroup.com